

Unit_8_Beer

MayaC

2024-02-27

R Markdown

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com> (<http://rmarkdown.rstudio.com>).

When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:

```
library(tidyverse)
```

```
## — Attaching core tidyverse packages — tidyverse 2.0.0 —
## ✓ dplyr      1.1.4      ✓ readr      2.1.4
## ✓ forcats    1.0.0      ✓ stringr    1.5.1
## ✓ ggplot2    3.4.4      ✓ tibble     3.2.1
## ✓ lubridate  1.9.3      ✓ tidyr      1.3.0
## ✓ purrr      1.0.2
## — Conflicts — tidyverse_conflicts() —
## ✖ dplyr::filter() masks stats::filter()
## ✖ dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(dplyr)
library(stringr)
library(ggplot2)
#connecting to the dataset – beer and breweries dataset provided for this project
beers <- read.csv("Beers.csv", header= TRUE)
breweries <- read.csv("Breweries.csv", header= TRUE)
#check data
head(beers)
```

```
##           Name Beer_ID  ABV IBU Brewery_id
## 1      Pub Beer   1436 0.050  NA      409
## 2    Devil's Cup   2265 0.066  NA      178
## 3 Rise of the Phoenix 2264 0.071  NA      178
## 4      Sinister   2263 0.090  NA      178
## 5    Sex and Candy 2262 0.075  NA      178
## 6    Black Exodus 2261 0.077  NA      178
##           Style Ounces
## 1    American Pale Lager      12
## 2    American Pale Ale (APA)    12
## 3          American IPA        12
## 4 American Double / Imperial IPA 12
## 5          American IPA        12
## 6          Oatmeal Stout      12
```

```
head(breweries)
```

```
## Brew_ID           Name           City State
## 1      1    NorthGate Brewing    Minneapolis  MN
## 2      2 Against the Grain Brewery    Louisville  KY
## 3      3  Jack's Abby Craft Lagers    Framingham  MA
## 4      4 Mike Hess Brewing Company    San Diego    CA
## 5      5  Fort Point Beer Company San Francisco  CA
## 6      6  COAST Brewing Company    Charleston  SC
```

```
#create data frames to answer questions
```

```
dfbeer = data.frame(BeerName = beers$Name, Beer_ID = beers$Beer_ID, ABV = beers$ABV, IBU = beers$IBU, Brew_ID = beers$Brewery_id, Style = beers$Style, Ounces = beers$Ounces)
```

```
dfbrewery = data.frame(Brew_ID = breweries$Brew_ID, BreweryName = breweries$Name, City = breweries$City, State = breweries$State)
```

```
#check data
```

```
head(dfbeer)
```

```
##           BeerName Beer_ID  ABV IBU Brew_ID           Style
## 1      Pub Beer   1436 0.050  NA      409    American Pale Lager
## 2    Devil's Cup   2265 0.066  NA      178    American Pale Ale (APA)
## 3 Rise of the Phoenix 2264 0.071  NA      178          American IPA
## 4      Sinister   2263 0.090  NA      178 American Double / Imperial IPA
## 5    Sex and Candy 2262 0.075  NA      178          American IPA
## 6    Black Exodus 2261 0.077  NA      178          Oatmeal Stout
## Ounces
## 1      12
## 2      12
## 3      12
## 4      12
## 5      12
## 6      12
```

```
head(dfbrewery)
```

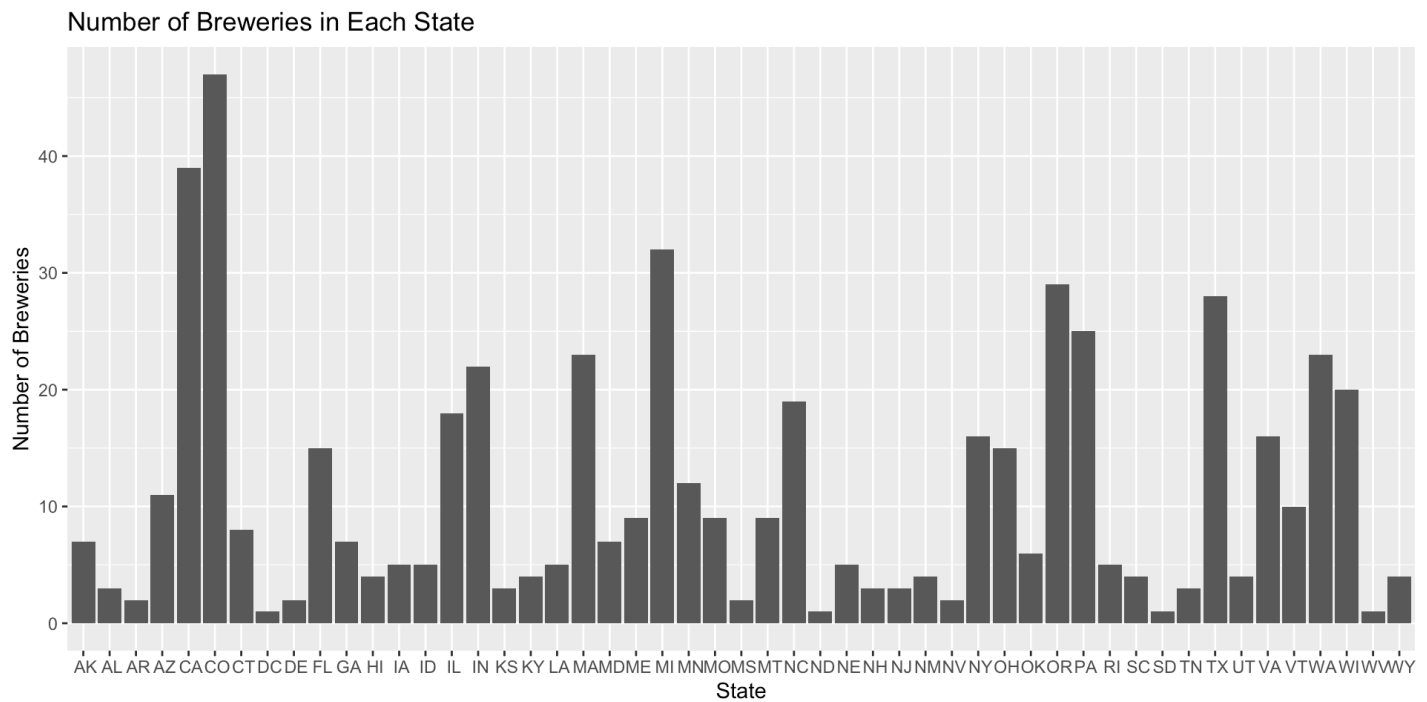
```
##   Brew_ID      BreweryName      City State
## 1      1      NorthGate Brewing  Minneapolis  MN
## 2      2  Against the Grain Brewery  Louisville  KY
## 3      3  Jack's Abby Craft Lagers  Framingham  MA
## 4      4  Mike Hess Brewing Company  San Diego  CA
## 5      5  Fort Point Beer Company  San Francisco  CA
## 6      6  COAST Brewing Company  Charleston  SC
```

```
#merge datasets -
dfcombined <- merge(dfbeer,dfbrewery,"Brew_ID")
```

```
#check data
head(dfcombined)
```

```
##   Brew_ID      BeerName Beer_ID  ABV IBU      Style
## 1      1  Get Together  2692 0.045  50      American IPA
## 2      1  Maggie's Leap  2691 0.049  26      Milk / Sweet Stout
## 3      1  Wall's End  2690 0.048  19      English Brown Ale
## 4      1  Pumphion  2689 0.060  38      Pumpkin Ale
## 5      1  Stronghold  2688 0.060  25      American Porter
## 6      1  Parapet ESB  2687 0.056  47  Extra Special / Strong Bitter (ESB)
##   Ounces      BreweryName      City State
## 1     16  NorthGate Brewing  Minneapolis  MN
## 2     16  NorthGate Brewing  Minneapolis  MN
## 3     16  NorthGate Brewing  Minneapolis  MN
## 4     16  NorthGate Brewing  Minneapolis  MN
## 5     16  NorthGate Brewing  Minneapolis  MN
## 6     16  NorthGate Brewing  Minneapolis  MN
```

```
#Question 1 - How many breweries are present in each state?
# Create a bar plot showing the count of breweries in each state
library(ggplot2)
ggplot(dfbrewery, aes(x = State)) +
  geom_bar() +
  labs(x = "State", y = "Number of Breweries", title = "Number of Breweries in Each State")
```



#Question 2 Merge beer data with the breweries data. Print the first 6 observations and the last six observations to check the merged file.

#first 6 observations

```
first_6_rows <- head(dfcombined,6)
print(first_6_rows)
```

```
##   Brew_ID   BeerName Beer_ID  ABV IBU                               Style
## 1      1  Get Together  2692 0.045  50                        American IPA
## 2      1 Maggie's Leap  2691 0.049  26                      Milk / Sweet Stout
## 3      1   Wall's End  2690 0.048  19                      English Brown Ale
## 4      1    Pumpion  2689 0.060  38                      Pumpkin Ale
## 5      1  Stronghold  2688 0.060  25                      American Porter
## 6      1  Parapet ESB  2687 0.056  47 Extra Special / Strong Bitter (ESB)
##   Ounces   BreweryName      City State
## 1    16 NorthGate Brewing  Minneapolis  MN
## 2    16 NorthGate Brewing  Minneapolis  MN
## 3    16 NorthGate Brewing  Minneapolis  MN
## 4    16 NorthGate Brewing  Minneapolis  MN
## 5    16 NorthGate Brewing  Minneapolis  MN
## 6    16 NorthGate Brewing  Minneapolis  MN
```

#last 6 observations

```
last_6_rows <- dfcombined[(nrow(dfcombined) - 5):nrow(dfcombined), ]
print(last_6_rows)
```

##	Brew_ID	BeerName	Beer_ID	ABV	IBU
## 2405	556	Pilsner Ukiah	98	0.055	NA
## 2406	557	Heinnieweisse Weissebier	52	0.049	NA
## 2407	557	Snapperhead IPA	51	0.068	NA
## 2408	557	Moo Thunder Stout	50	0.049	NA
## 2409	557	Porkslap Pale Ale	49	0.043	NA
## 2410	558	Urban Wilderness Pale Ale	30	0.049	NA
##	Style	Ounces	BreweryName	City	
## 2405	German Pilsener	12	Ukiah Brewing Company	Ukiah	
## 2406	Hefeweizen	12	Butternuts Beer and Ale	Garrattsville	
## 2407	American IPA	12	Butternuts Beer and Ale	Garrattsville	
## 2408	Milk / Sweet Stout	12	Butternuts Beer and Ale	Garrattsville	
## 2409	American Pale Ale (APA)	12	Butternuts Beer and Ale	Garrattsville	
## 2410	English Pale Ale	12	Sleeping Lady Brewing Company	Anchorage	
##	State				
## 2405	CA				
## 2406	NY				
## 2407	NY				
## 2408	NY				
## 2409	NY				
## 2410	AK				

#managing missing values – ABV values are missing for 62 of the 2,405 rows and IBU values are missing for 1,005 of 2,410 rows. The ABV values look like there are missing completely at random. While I can't find any evidence of "missing at random" or "not missing at random" for IBU values, my gut says I am not looking at all the angles. I made a call to calculate the mean and populate the missing data with mean values. This has been done for both ABV and IBU. All the below questions are answered based on this dataset.

calculate mean for both variables at the Style level

library(dplyr)

```
dfcombined_Mean <- dfcombined %>%
```

```
  group_by(Style) %>%
```

```
  summarise(mean_ABV = mean(ABV, na.rm = TRUE),
```

```
            mean_IBU = mean(IBU, na.rm = TRUE))
```

#populate NA values for ABV and IBU with mean values

Merge dfcombined with dfcombined_Mean to get mean values for each style

```
dfcombined_updated <- dfcombined %>%
```

```
  left_join(dfcombined_Mean, by = "Style")
```

Replace NA values in ABV and IBU columns with corresponding mean values

```
dfcombined_updated$ABV[is.na(dfcombined_updated$ABV)] <- dfcombined_updated$mean_ABV[is.na(dfcombined_updated$ABV)]
```

```
dfcombined_updated$IBU[is.na(dfcombined_updated$IBU)] <- dfcombined_updated$mean_IBU[is.na(dfcombined_updated$IBU)]
```

Remove unnecessary columns (mean_ABV, mean_IBU)

```
dfcombined_updated <- dfcombined_updated %>%
```

```
  select(-mean_ABV, -mean_IBU)
```

Print rows where IBU or ABV is NA

```
dim(dfcombined_updated %>% filter(is.na(IBU) | is.na(ABV)))
```

```
## [1] 52 10
```

there are 52 rows where none of the IBU values for that Style was populated. Hence we couldn't replace the NA with mean values. These were removed from the dataset for the below questions.

#Question 4 – Compute the median alcohol content and international bitterness unit for each state. Plot a bar chart to compare.

Remove NA values from ABV and IBU columns

```
dfcombined_clean <- dfcombined[!is.na(dfcombined$ABV) & !is.na(dfcombined$IBU), ]
```

```
dfcombined_updated_clean <- dfcombined_updated[!is.na(dfcombined_updated$ABV) & !is.na(dfcombined_updated$IBU), ]
```

#check data

```
dim(dfcombined)
```

```
## [1] 2410 10
```

```
dim(dfcombined_clean)
```

```
## [1] 1405 10
```

```
dim(dfcombined_updated)
```

```
## [1] 2410 10
```

```
dim(dfcombined_updated_clean)
```

```
## [1] 2358 10
```

```
# Calculate median ABV and median IBU for each state
medians_by_state <- dfcombined_updated_clean %>%
  group_by(State) %>%
  summarise(Median_ABV = median(ABV), Median_IBU = median(IBU, na.rm = TRUE))
Overall_Median <- median(medians_by_state$Median_ABV)
Overall_Median_I <- median(medians_by_state$Median_IBU)

# check data
head(medians_by_state,10)
```

```
## # A tibble: 10 × 3
##   State Median_ABV Median_IBU
##   <chr>      <dbl>      <dbl>
## 1 " AK"      0.056        33.8
## 2 " AL"      0.06         39.5
## 3 " AR"      0.052        36.3
## 4 " AZ"      0.0575       22.2
## 5 " CA"      0.058         40
## 6 " CO"      0.06         36.3
## 7 " CT"      0.06         36.3
## 8 " DC"      0.0625       28.1
## 9 " DE"      0.0598       59.8
## 10 " FL"     0.0555       35.6
```

```


#plot the median values
bar1 <- ggplot(medians_by_state, aes(x = State, y = Median_ABV)) + # Move aes() inside
ggplot()
  geom_col(fill = 'blue') +
  labs(x = "State", y = "Median ABV", title = "Median Alcohol by Volume of Beers in Each
State") +
  theme(
    plot.title = element_text(size = 20),
    axis.text.x = element_text(angle = 45, hjust = 1),
    axis.text.y = element_text(size = 12),
    axis.title = element_text(size = 14)
  ) +
  geom_hline(yintercept = Overall_Median, color = "red") # Add a horizontal line for ov
erall median

bar2 = ggplot(medians_by_state, aes(x = State, y = Median_IBU)) +
  geom_col(fill = 'green') +
  labs(x = "State", y = "Median IBU", title = "Median International Bitterness Unit of B
eers in Each State") +
  theme(
    plot.title = element_text(size = 20),
    axis.text.x = element_text(angle = 45, hjust = 1),
    axis.text.y = element_text(size = 12),
    axis.title = element_text(size = 14)
  ) +
  geom_hline(yintercept = Overall_Median_I, color = "red") # Add a horizontal line for
overall median


#display both plots
library(patchwork)
bar1 / bar2

```

Median Alcohol by Volume of Beers in Each State



Median International Bitterness Unit of Beers in Each State



#Question 5 – Which state has the maximum alcoholic (ABV) beer? Which state has the most bitter (IBU) beer?

#state with maximum alcoholic beer

```
max_abv <- max(dfcombined_updated_clean$ABV, na.rm = TRUE)
max_abv_index <- which.max(dfcombined_updated_clean$ABV)
state_with_max_abv <- dfcombined_updated_clean$State[max_abv_index]
print(state_with_max_abv)
```

```
## [1] " CO"
```

```
cat("State with maximum alcohol by volume beer is :",state_with_max_abv, "with ABV of",max_abv,"\n")
```

```
## State with maximum alcohol by volume beer is : CO with ABV of 0.128
```

#state with most bitter beer

```
max_ibu <- max(dfcombined_updated_clean$IBU,na.rm = TRUE)
max_ibu_index <- which.max(dfcombined_updated_clean$IBU)
state_with_max_ibu <- dfcombined_updated_clean$State[max_ibu_index]
print(state_with_max_ibu)
```

```
## [1] " OR"
```

```
cat("State with most bitter beer is :",state_with_max_ibu, "with IBU of",max_ibu,"\n")
```

```
## State with most bitter beer is : OR with IBU of 138
```

```

#Question 6 - Comment on the summary statistics and distribution of the ABV variable.
#create a new region column
midwest <- c("IA","IL", "IN", "KS", "MI", "MN", "MO", "ND","NE", "OH", "SD", "WI")
northeast <- c("CT","MA", "ME", "NH", "NJ", "NY", "PA", "RI", "VT")
south <- c("AL","AR","DC","DE","FL","GA","KY","LA","MD","MS","NC","OK","SC","TN","TX","V
A","WV")
west <- c("AK","AZ","CA","CO","HI","ID","MT","NM","NV","OR","UT","WA","WY")

# Clean the State column to remove leading and trailing whitespace
dfcombined_updated_clean$State <- trimws(dfcombined_updated_clean$State)

# Create a new column named "Region" and assign the region based on state
dfcombined_updated_clean$Region <- ifelse(dfcombined_updated_clean$State %in% midwest,
"Midwest",
                                     ifelse(dfcombined_updated_clean$State %in% northeast,
"Northeast",
                                     ifelse(dfcombined_updated_clean$State %in% sout
h, "South",
                                     ifelse(dfcombined_updated_clean$State %i
n% west, "West", "Unknown"))))

# Calculate summary statistics
ABV <- dfcombined_updated_clean$ABV
mean_ABV <- mean(ABV, na.rm = TRUE) # Mean
median_ABV <- median(ABV, na.rm = TRUE) # Median
min_ABV <- min(ABV, na.rm = TRUE) # Minimum
max_ABV <- max(ABV, na.rm = TRUE) # Maximum
q1_ABV <- quantile(ABV, probs = 0.25, na.rm = TRUE) # First quartile (25th percentile)
q3_ABV <- quantile(ABV, probs = 0.75, na.rm = TRUE) # Third quartile (75th percentile)

# Print summary statistics
cat("Mean ABV:", mean_ABV, "\n")

```

```
## Mean ABV: 0.05972379
```

```
cat("Median ABV:", median_ABV, "\n")
```

```
## Median ABV: 0.056
```

```
cat("Minimum ABV:", min_ABV, "\n")
```

```
## Minimum ABV: 0.027
```

```
cat("Maximum ABV:", max_ABV, "\n")
```

```
## Maximum ABV: 0.128
```

```
cat("First Quartile ABV:", q1_ABV, "\n")
```


```
## First Quartile ABV: 0.05
```

```
cat("Third Quartile ABV:", q3_ABV, "\n")
```

```
## Third Quartile ABV: 0.067
```

```
#distribution
ggplot(dfcombined_updated_clean,aes(x=ABV)) +
  geom_histogram()
```


```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
#using another way
hist(ABV, main = "Distribution of Alcohol By Volume (ABV)", xlab = "ABV", ylab = "Frequency")
```


```
# Add summary statistics to the plot
abline(v = mean_ABV, col = "red", lwd = 2) # Mean line
abline(v = median_ABV, col = "blue", lwd = 2) # Median line
legend("topright", legend = c("Mean", "Median"), col = c("red", "blue"), lwd = 2)
```

Distribution of Alcohol By Volume (ABV)



```
#boxplot
# Box plot with color and labels
ggplot(dfcombined_updated_clean, aes(x = "", y = ABV, fill = Region)) +
  geom_boxplot() +
  labs(title = "Distribution of Alcohol By Volume (ABV) by Region", x = NULL, y = "Alcohol By Volume (ABV)", fill = "Region")
```

Distribution of Alcohol By Volume (ABV) by Region



#summary statement

"The summary statistics for the ABV variable show that the mean alcohol by volume is 5.9%, with a median of 5.7%. The range of ABV values extends from 2.7% to 12.5%, with 5% as the first quartile and 6.8% as the third quartile. The distribution of ABV values appears to be right-skewed, indicating that the majority of beers have lower alcohol content. However, there are some outliers with exceptionally high ABV values, which may represent specialty beers."

```
## [1] "The summary statistics for the ABV variable show that the mean alcohol by volume is 5.9%, with a median of 5.7%. The range of ABV values extends from 2.7% to 12.5%, with 5% as the first quartile and 6.8% as the third quartile. The distribution of ABV values appears to be right-skewed, indicating that the majority of beers have lower alcohol content. However, there are some outliers with exceptionally high ABV values, which may represent specialty beers."
```

#question 7 Is there an apparent relationship between the bitterness of the beer and its alcoholic content? Draw a scatter plot

#create a new region column

```
midwest <- c("IA","IL","IN","KS","MI","MN","MO","ND","NE","OH","SD","WI")
northeast <- c("CT","MA","ME","NH","NJ","NY","PA","RI","VT")
south <- c("AL","AR","DC","DE","FL","GA","KY","LA","MD","MS","NC","OK","SC","TN","TX","VA","WV")
west <- c("AK","AZ","CA","CO","HI","ID","MT","NM","NV","OR","UT","WA","WY")
```

Clean the State column to remove leading and trailing whitespace

```
dfcombined_updated_clean$State <- trimws(dfcombined_updated_clean$State)
```

Create a new column named "Region" and assign the region based on state

```
dfcombined_updated_clean$Region <- ifelse(dfcombined_updated_clean$State %in% midwest,
"Midwest",
                                     ifelse(dfcombined_updated_clean$State %in% northeast,
"Northeast",
                                     ifelse(dfcombined_updated_clean$State %in% south, "South",
                                     ifelse(dfcombined_updated_clean$State %in% west, "West", "Unknown"))))
```

Check the data

```
unique(dfcombined_updated_clean$Region)
```

```
## [1] "Midwest" "South" "Northeast" "West"
```

```
head(dfcombined_updated_clean)
```

##	Brew_ID	BeerName	Beer_ID	ABV	IBU	Style
## 1	1	Get Together	2692	0.045	50	American IPA
## 2	1	Maggie's Leap	2691	0.049	26	Milk / Sweet Stout
## 3	1	Wall's End	2690	0.048	19	English Brown Ale
## 4	1	Pumpkin	2689	0.060	38	Pumpkin Ale
## 5	1	Stronghold	2688	0.060	25	American Porter
## 6	1	Parapet ESB	2687	0.056	47	Extra Special / Strong Bitter (ESB)


##	Ounces	BreweryName	City	State	Region
## 1	16	NorthGate Brewing	Minneapolis	MN	Midwest
## 2	16	NorthGate Brewing	Minneapolis	MN	Midwest
## 3	16	NorthGate Brewing	Minneapolis	MN	Midwest
## 4	16	NorthGate Brewing	Minneapolis	MN	Midwest
## 5	16	NorthGate Brewing	Minneapolis	MN	Midwest
## 6	16	NorthGate Brewing	Minneapolis	MN	Midwest

```
#draw scatterplot
```

```
dfcombined_updated_clean %>% select(State,Region,ABV,IBU) %>%
```

```
ggplot(aes(x= ABV, y= IBU,color = Region))+
  geom_point()+labs(title = "Scatterplot of Alcohol by Volume Vs. International Bitterness Unit", x= "Alcohol by Volume", y= "International Bitterness Unit")
```

Scatterplot of Alcohol by Volume Vs. International Bitterness Unit



```
#summary statement
```

"It looks like there is a slight positive association between Alcohol by Volume and International Bitterness Unit. While there is a general tendency for the bitterness to move in the same direction as alcohol by volume, there is also considerable variability around the trend line. Beers from all regions are represented in the variability around the trend line."

```
## [1] "It looks like there is a slight positive association between Alcohol by Volume and International Bitterness Unit. While there is a general tendency for the bitterness to move in the same direction as alcohol by volume, there is also considerable variability around the trend line. Beers from all regions are represented in the variability around the trend line."
```

```
#histogram by region for both IBU and ABV values
```

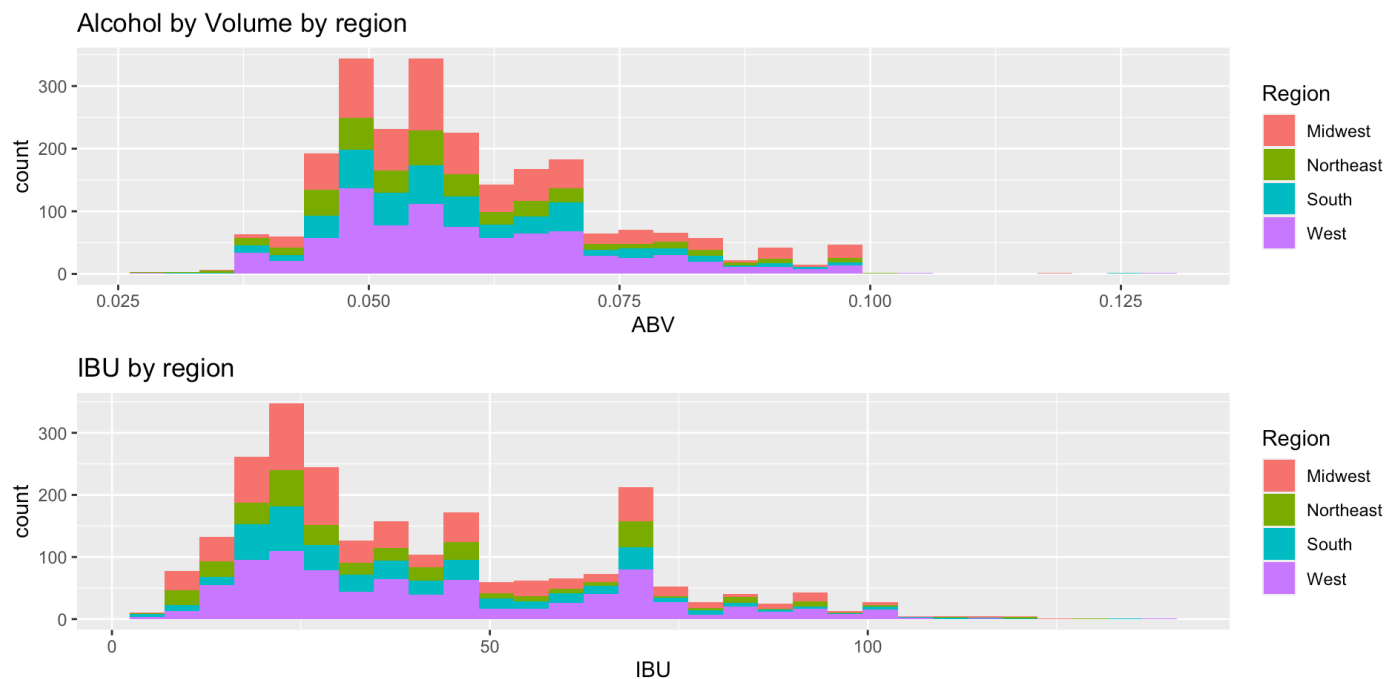
```
ABV <- dfcombined_updated_clean %>%
  select(State, Region, ABV, IBU) %>%
  ggplot(aes(x = ABV, fill = Region)) +
  geom_histogram() +
  labs(title = "Alcohol by Volume by region")
```

```
IBU <- dfcombined_updated_clean %>%
  select(State, Region, ABV, IBU) %>%
  ggplot(aes(x = IBU, fill = Region)) +
  geom_histogram() +
  labs(title = "IBU by region")
```

```
library(patchwork)
```


```
ABV / IBU
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
#histogram by region for ABV and IBU values in facet wrap style
dfcombined_updated_clean %>%
  select(State, Region, ABV, IBU) %>%
  ggplot(aes(x = ABV, fill = Region)) +
  geom_histogram() +
  labs(title = "Alcohol by Volume by region") +
  facet_wrap(~Region)
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
dfcombined_updated_clean %>%
  select(State, Region, ABV, IBU) %>%
  ggplot(aes(x = IBU, fill = Region)) +
  geom_histogram() +
  labs(title = "IBU by region") +
  facet_wrap(~Region)
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

#question 9 – use knn to explore relationship of ABV & IBU between IPA and other Ales
#creating a data frame with styles of IPA and "ale" in the name

```
library(dplyr)
library(stringr)
```

```
# Filter rows where the Style column contains "IPA" or "Ale"
dfIPA_Ale <- dfcombined_updated_clean %>%
  filter(str_detect(Style, "IPA") | grepl("\\bAle\\b", Style))
```

```
# View the filtered data
head(dfIPA_Ale)
```

##	Brew_ID	BeerName	Beer_ID	ABV	IBU	Style
## 1	1	Get Together	2692	0.045	50	American IPA
## 2	1	Wall's End	2690	0.048	19	English Brown Ale
## 3	1	Pumpkin	2689	0.060	38	Pumpkin Ale
## 4	2	Citra Ass Down	2686	0.080	68	American Double / Imperial IPA
## 5	2	A Beer	2683	0.042	42	American Pale Ale (APA)
## 6	2	Flesh Gourd'n	2681	0.066	21	Pumpkin Ale

##	Ounces	BreweryName	City	State	Region
## 1	16	NorthGate Brewing	Minneapolis	MN	Midwest
## 2	16	NorthGate Brewing	Minneapolis	MN	Midwest
## 3	16	NorthGate Brewing	Minneapolis	MN	Midwest
## 4	16	Against the Grain Brewery	Louisville	KY	South
## 5	16	Against the Grain Brewery	Louisville	KY	South
## 6	16	Against the Grain Brewery	Louisville	KY	South

```
dim(dfIPA_Ale)
```

```
## [1] 1533 11
```

```
dim(dfcombined_updated_clean)
```

```
## [1] 2358 11
```


```
# Add a Type column to differentiate IPA from other Ales
dfIPA_Ale$Type <- ifelse(str_detect(dfIPA_Ale$Style, "IPA"), "IPA", "Ale")
```

```
#check data
head(dfIPA_Ale)
```

```
## Brew_ID BeerName Beer_ID ABV IBU Style
## 1 1 Get Together 2692 0.045 50 American IPA
## 2 1 Wall's End 2690 0.048 19 English Brown Ale
## 3 1 Pumpion 2689 0.060 38 Pumpkin Ale
## 4 2 Citra Ass Down 2686 0.080 68 American Double / Imperial IPA
## 5 2 A Beer 2683 0.042 42 American Pale Ale (APA)
## 6 2 Flesh Gourd'n 2681 0.066 21 Pumpkin Ale
## Ounces BreweryName City State Region Type
## 1 16 NorthGate Brewing Minneapolis MN Midwest IPA
## 2 16 NorthGate Brewing Minneapolis MN Midwest Ale
## 3 16 NorthGate Brewing Minneapolis MN Midwest Ale
## 4 16 Against the Grain Brewery Louisville KY South IPA
## 5 16 Against the Grain Brewery Louisville KY South Ale
## 6 16 Against the Grain Brewery Louisville KY South Ale
```


```
# draw scatterplot
dfIPA_Ale %>%
  select(State, Region, Type, ABV, IBU) %>%
  ggplot(aes(x = ABV, y = IBU, color = Type)) +
  geom_point() +
  labs(
    title = "Scatterplot of Alcohol by Volume Vs. International Bitterness Unit",
    x = "Alcohol by Volume",
    y = "International Bitterness Unit"
  )
```

Scatterplot of Alcohol by Volume Vs. International Bitterness Unit



```
# draw scatterplot by region
dfIPA_Ale %>%
  select(State, Region, Type, ABV, IBU) %>%
  ggplot(aes(x = ABV, y = IBU, color = Type)) +
  geom_point() +
  labs(
    title = "Scatterplot of Alcohol by Volume Vs. International Bitterness Unit",
    x = "Alcohol by Volume",
    y = "International Bitterness Unit"
  ) +
  facet_wrap(~Region)
```

Scatterplot of Alcohol by Volume Vs. International Bitterness Unit



```
#training and test datasets - knn classification
```

```
library(class)
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
##
```

```
## Attaching package: 'caret'
```

```
##
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
## lift
```

```
library(e1071)
```

```
library(dplyr)
```

```
set.seed(6)
```

```
splitPerc = .7
```

```
trainIndices = sample(1:dim(dfIPA_Ale)[1],round(splitPerc * dim(dfIPA_Ale)[1]))
```

```
train = dfIPA_Ale[trainIndices,]
```

```
test = dfIPA_Ale[-trainIndices,]
```

```
#check data
```

```
nrow(train)
```

```
## [1] 1073
```

```
nrow(test)
```

```
## [1] 460
```

```
nrow(dfIPA_Ale)
```

```
## [1] 1533
```

```
# k = 5
```

```
classifications = knn(train[,c(4,5)],test[,c(4,5)],train$Type, prob = TRUE, k = 7)
```

```
table(classifications,test$Type)
```

```
##
```

```
## classifications Ale IPA
```

```
## Ale 256 35
```

```
## IPA 22 147
```

```
confusionMatrix(table(classifications,test$Type), mode = "everything")
```

```
## Confusion Matrix and Statistics
##
##
## classifications Ale IPA
##           Ale 256  35
##           IPA  22 147
##
##           Accuracy : 0.8761
##           95% CI : (0.8425, 0.9048)
##           No Information Rate : 0.6043
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.7377
##
## Mcnemar's Test P-Value : 0.112
##
##           Sensitivity : 0.9209
##           Specificity : 0.8077
##           Pos Pred Value : 0.8797
##           Neg Pred Value : 0.8698
##           Precision : 0.8797
##           Recall : 0.9209
##           F1 : 0.8998
##           Prevalence : 0.6043
##           Detection Rate : 0.5565
##           Detection Prevalence : 0.6326
##           Balanced Accuracy : 0.8643
##
##           'Positive' Class : Ale
##
```

```
#trying naive bayes
#Train a NB model based on the training set using ABV and IBU.
model = naiveBayes(train[,c(4,5)],train$Type)
#Use the model to predict the Type in the test set and use those results to evaluate the
model based on accuracy, sensitivity and specificity.
table(predict(model,test[,c(4,5)]),test$Type)
```

```
##
##           Ale IPA
##           Ale 243  34
##           IPA  35 148
```

```
confusionMatrix(table(predict(model,test[,c(4,5)]),test$Type), mode = "everything")
```

```
## Confusion Matrix and Statistics
##
##
##      Ale IPA
## Ale 243  34
## IPA  35 148
##
##              Accuracy : 0.85
##              95% CI : (0.814, 0.8814)
##      No Information Rate : 0.6043
##      P-Value [Acc > NIR] : <2e-16
##
##              Kappa : 0.6866
##
## Mcnemar's Test P-Value : 1
##
##      Sensitivity : 0.8741
##      Specificity : 0.8132
##      Pos Pred Value : 0.8773
##      Neg Pred Value : 0.8087
##      Precision : 0.8773
##      Recall : 0.8741
##      F1 : 0.8757
##      Prevalence : 0.6043
##      Detection Rate : 0.5283
##      Detection Prevalence : 0.6022
##      Balanced Accuracy : 0.8436
##
##      'Positive' Class : Ale
##
```

```
#bonus question - additional insights from the data
# Load necessary packages
library(ggplot2)
library(dplyr)
library(maps)
```

```
##
## Attaching package: 'maps'
##
## The following object is masked from 'package:purrr':
##
##      map
```

```

# Create a mapping table for state abbreviations to full names
state_mapping <- data.frame(StateAbbrev = c("AL", "AK", "AZ", "AR", "CA", "CO", "CT", "D
E", "FL", "GA",
                                         "HI", "ID", "IL", "IN", "IA", "KS", "KY", "L
A", "ME", "MD",
                                         "MA", "MI", "MN", "MS", "MO", "MT", "NE", "N
V", "NH", "NJ",
                                         "NM", "NY", "NC", "ND", "OH", "OK", "OR", "P
A", "RI", "SC",
                                         "SD", "TN", "TX", "UT", "VT", "VA", "WA", "W
V", "WI", "WY"),
                           StateFull = c("alabama", "alaska", "arizona", "arkansas", "ca
lifornia", "colorado",
                                         "connecticut", "delaware", "florida", "georgi
a", "hawaii", "idaho",
                                         "illinois", "indiana", "iowa", "kansas", "kentu
cky", "louisiana",
                                         "maine", "maryland", "massachusetts", "michiga
n", "minnesota",
                                         "mississippi", "missouri", "montana", "nebrask
a", "nevada",
                                         "new hampshire", "new jersey", "new mexico", "n
ew york", "north carolina",
                                         "pennsylvania",
                                         "rhode island", "south carolina", "south dakot
a", "tennessee",
                                         "texas", "utah", "vermont", "virginia", "washin
gton", "west virginia",
                                         "wisconsin", "wyoming"))

# Merge the mapping table with your data frame to get full state names
dfcombined_updated_clean <- merge(dfcombined_updated_clean, state_mapping, by.x = "Stat
e", by.y = "StateAbbrev")

# Aggregate the number of beers available by state
state_beer_data <- dfcombined_updated_clean %>%
  group_by(StateFull) %>%
  summarise(TotalBeersAvailable = n())

# Load US map data
us_map <- map_data("state")

# Merge your aggregated data with map data
merged_data <- merge(us_map, state_beer_data, by.x = "region", by.y = "StateFull", all.x
= TRUE)

# Plot the map
ggplot() +
  geom_map(data = merged_data, map = merged_data,
           aes(x = long, y = lat, map_id = region, fill = TotalBeersAvailable),
           color = "black", size = 0.15) +


```

```
scale_fill_gradient(low = "lightblue", high = "darkblue", na.value = "grey") + # Adjust color gradient
labs(title = "Number of beer choices by State", fill = "Total Beers Available") + # Label legend
coord_fixed(1.3) + # Adjust aspect ratio
theme_void() # Remove axis and grid lines
```

```
## Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
## i Please use `linewidth` instead.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.
```

```
## Warning in geom_map(data = merged_data, map = merged_data, aes(x = long, :
## Ignoring unknown aesthetics: x and y
```

Number of beer choices by State



```
#bonus question - most popular ounces
#plot bar
# Aggregate the number of beers available by Ounces
Ounce_data <- dfcombined_updated_clean %>%
  group_by(Ounces) %>%
  summarise(TotalBeersAvailable = n())

# Plot the bar graph
ggplot(Ounce_data, aes(x = Ounces, y = TotalBeersAvailable)) +
  geom_col(show.legend = FALSE, fill = "brown") +
  labs(title = "Number of beers by ounce packaging")
```